

```
In [3]: import pandas as pd
import matplotlib.pyplot as plt
import numpy as np
from collections import Counter
from sklearn.preprocessing import StandardScaler

import imblearn as il
from imblearn.over_sampling import SMOTE
from imblearn.combine import SMOTEENN
from sklearn.model_selection import train_test_split
%matplotlib inline

from sklearn.linear_model import LogisticRegression
from sklearn.ensemble import RandomForestClassifier
from sklearn.metrics import recall_score
from sklearn.metrics import average_precision_score
from sklearn.metrics import precision_recall_curve

pd.set_option('display.max_rows', 500)
pd.set_option('display.max_columns', 500)
pd.set_option('display.float_format', lambda x: '%.4f' % x)
```

```
In [2]: data = pd.read_csv("/Users/sonia/Downloads/creditcard.csv")
```


```
In [4]: count_classes = pd.value_counts(data['Class'], sort = True).sort_index()

count_classes
```

```
Out[4]: 0    284315
1         492
Name: Class, dtype: int64
```

```
In [5]: plt.figure(figsize=(10,5))
count_classes.plot(kind = 'bar')
plt.title("Fraud class histogram")
plt.xlabel("Class")
plt.ylabel("Frequency")
```

```
Out[5]: Text(0,0.5,'Frequency')
```



```
In [6]: data.isnull().sum()
```

```
Out[6]: Time      0
        V1        0
        V2        0
        V3        0
        V4        0
        V5        0
        V6        0
        V7        0
        V8        0
        V9        0
        V10       0
        V11       0
        V12       0
        V13       0
        V14       0
        V15       0
        V16       0
        V17       0
        V18       0
        V19       0
        V20       0
        V21       0
        V22       0
        V23       0
        V24       0
        V25       0
        V26       0
        V27       0
        V28       0
        Amount    0
        Class     0
        dtype: int64
```

```
In [7]: data.describe()
```

```
Out[7]:
```

	Time	V1	V2	V3	V4	V5	V6
count	284807.0000	284807.0000	284807.0000	284807.0000	284807.0000	284807.0000	284807.0000
mean	94813.8596	0.0000	0.0000	-0.0000	0.0000	-0.0000	0.0000
std	47488.1460	1.9587	1.6513	1.5163	1.4159	1.3802	1.3302
min	0.0000	-56.4075	-72.7157	-48.3256	-5.6832	-113.7433	-26.1600
25%	54201.5000	-0.9204	-0.5985	-0.8904	-0.8486	-0.6916	-0.7600
50%	84692.0000	0.0181	0.0655	0.1798	-0.0198	-0.0543	-0.2700
75%	139320.5000	1.3156	0.8037	1.0272	0.7433	0.6119	0.3900
max	172792.0000	2.4549	22.0577	9.3826	16.8753	34.8017	73.3000

```
In [9]: cols = list(data.columns.values)
```

```
In [10]: cols
```

```
Out[10]: ['Time',  
          'V1',  
          'V2',  
          'V3',  
          'V4',  
          'V5',  
          'V6',  
          'V7',  
          'V8',  
          'V9',  
          'V10',  
          'V11',  
          'V12',  
          'V13',  
          'V14',  
          'V15',  
          'V16',  
          'V17',  
          'V18',  
          'V19',  
          'V20',  
          'V21',  
          'V22',  
          'V23',  
          'V24',  
          'V25',  
          'V26',  
          'V27',  
          'V28',  
          'Amount',  
          'Class']
```

```
In [11]: X = data.iloc[:, data.columns != 'Class']  
         Y = data.iloc[:, data.columns == 'Class']
```

```
In [12]: Y.shape
```

```
Out[12]: (284807, 1)
```

```
In [13]: X.shape
```

```
Out[13]: (284807, 30)
```

```
In [14]: print('Class labels:', np.unique(Y))
```

```
Class labels: [0 1]
```

```
In [15]: X_train, X_test, Y_train, Y_test = train_test_split(X, Y, test_size=0.1,
```

```
In [16]: train, x_val, y_train, y_val = train_test_split(X_train, Y_train, test_size
```

```
In [17]: lr = LogisticRegression(random_state=0)
```

```
In [18]: lr.fit(x_train, y_train)
```

```
/anaconda3/lib/python3.6/site-packages/sklearn/linear_model/logistic.py:433: FutureWarning: Default solver will be changed to 'lbfgs' in 0.22. Specify a solver to silence this warning.
```

```
FutureWarning)
```

```
/anaconda3/lib/python3.6/site-packages/sklearn/utils/validation.py:761: DataConversionWarning: A column-vector y was passed when a 1d array was expected. Please change the shape of y to (n_samples, ), for example using ravel().
```

```
y = column_or_1d(y, warn=True)
```

```
Out[18]: LogisticRegression(C=1.0, class_weight=None, dual=False, fit_intercept=True,
```

```
intercept_scaling=1, max_iter=100, multi_class='warn',  
n_jobs=None, penalty='l2', random_state=0, solver='warn',  
tol=0.0001, verbose=0, warm_start=False)
```

```
In [19]: print('Validation Results')  
print('Accuracy: ', lr.score(x_val, y_val))  
print('Recall Score: ', recall_score(y_val, lr.predict(x_val)))
```

```
Validation Results
```

```
Accuracy: 0.9990246947294503
```

```
Recall Score: 0.5121951219512195
```

```
In [20]: print('Test Results')  
print('Accuracy: ', lr.score(X_test, Y_test))  
print('Recall Score: ', recall_score(Y_test, lr.predict(X_test)))
```

```
Test Results
```

```
Accuracy: 0.9990168884519505
```

```
Recall Score: 0.5636363636363636
```

```
In [21]: sm = SMOTE(random_state=12, ratio = 'auto', kind = 'regular')
```

```
In [22]: x_train_res, y_train_res = sm.fit_sample(x_train, y_train)

print('Resampled dataset shape {}'.format(Counter(y_train_res)))
```

```
/anaconda3/lib/python3.6/site-packages/sklearn/utils/validation.py:761
: DataConversionWarning: A column-vector y was passed when a 1d array
was expected. Please change the shape of y to (n_samples, ), for examp
le using ravel().
```

```
y = column_or_1d(y, warn=True)
```

```
Resampled dataset shape Counter({0: 230297, 1: 230297})
```

```
In [23]: sm_logr = LogisticRegression(random_state=0)

sm_logr.fit(x_train_res, y_train_res)
```

```
/anaconda3/lib/python3.6/site-packages/sklearn/linear_model/logistic.p
y:433: FutureWarning: Default solver will be changed to 'lbfgs' in 0.2
2. Specify a solver to silence this warning.
```

```
FutureWarning)
```

```
Out[23]: LogisticRegression(C=1.0, class_weight=None, dual=False, fit_intercept
=True,

        intercept_scaling=1, max_iter=100, multi_class='warn',
        n_jobs=None, penalty='l2', random_state=0, solver='warn',
        tol=0.0001, verbose=0, warm_start=False)
```

```
In [24]: print('Validation Results')
print('Accuracy: ', sm_logr.score(x_val, y_val))
print('Recall Score: ', recall_score(y_val, sm_logr.predict(x_val)))
```

```
Validation Results
```

```
Accuracy: 0.9855264697850428
```

```
Recall Score: 0.7804878048780488
```

```
In [25]: print('Test Results')
print('Accuracy: ', sm_logr.score(X_test, Y_test))
print('Recall Score: ', recall_score(Y_test, sm_logr.predict(X_test)))
```

```
Test Results
```

```
Accuracy: 0.9853937712861206
```

```
Recall Score: 0.9090909090909091
```

```
In [ ]:
```

